

Results

Results I

This section presents the experimental results from the gradient descent optimization study, including convergence analysis and performance comparisons. Every table, figure, and quantitative assertion in this section was compiled autonomously by the template's `infrastructure.reporting` subsystem executing the optimization analysis script. No manual transcription was permitted.


Convergence Analysis I

Convergence Analysis II

Convergence Trajectories

[@fig:convergence] illustrates the convergence behavior of gradient descent for different step sizes, starting from the initial point $x_0 = 0$. The algorithm iteratively updates the solution using the rule $x_{k+1} = x_k - \alpha \nabla f(x_k)$.

Gradient Descent Convergence Analysis
Quadratic Minimization



Quantitative Results I

The optimization results for different step sizes are synthesized computationally by orchestrating `infrastructure.reporting.executive_reporter`, feeding directly into the optimization results data that acts as the source of truth for `[@tbl:opt_results]`. Rows follow `experiment.step_sizes` in `config.yaml`; the body rows below are injected at render time from that CSV (`RESULT_TABLE_ROWS` in `scripts/z_generate_manuscript_variables.py`).

Table 1: Gradient descent outcomes per configured step size: state at termination, iteration count capped by $N_{\max} = 1000$, and the converged flag from `gradient_descent()` using $\|\nabla f\| < 10^{-8}$. Rows marked “No” either hit the iteration cap before meeting the gradient tolerance (small α) or correspond to unstable dynamics when $|1 - \alpha| \geq 1$.

Step Size ()	Final Solution	Objective Value	Iterations	Converged
0.01	1.0000	-0.5000	1000	No
0.10	1.0000	-0.5000	175	Yes
0.50	1.0000	-0.5000	27	Yes
1.00	1.0000	-0.5000	1	Yes

Quantitative Results II

1.50
2.50

1.0000
-

-0.5000
inf

27
1000

Yes
No

1233840596906177455240913358131801868905057899753059079921866755540991000896055603956


Convergence Rate Analysis I

Convergence Rate Analysis II

Theoretical vs Empirical Convergence

Modern convergence analysis builds on foundational work in gradient methods [nesterov2013gradient].

[fig:convergence_rate] provides a comparative analysis of convergence rates across different step sizes, validating theoretical predictions against empirical results.



Performance Analysis I

Convergence Speed

The results show a clear trade-off between step size and convergence speed:

- ▶ Small step sizes require more iterations but provide stable convergence
- ▶ Large step sizes converge faster but may be less stable in more complex problems

Performance Analysis II

Solution Accuracy

4 of 6 tested step sizes achieved the analytical optimum within numerical precision:

- ▶ Target solution: $x = 1.0$ (relative error $< 10^{-4}$ for converged settings)
- ▶ Target objective: $f(x) = -0.5$ (absolute error $< 10^{-8}$ for converged settings)

Divergent step sizes ($\alpha \geq 2$) confirm the theoretical instability boundary, demonstrating that gradient descent with fixed step size requires $\alpha < 2/\lambda_{\max}$ for convergence on quadratic objectives.

Algorithm Characteristics I

Strengths

- ▶ **Simplicity:** Easy to implement and understand
- ▶ **Generality:** Applicable to any differentiable objective function
- ▶ **Reliability:** Converges for convex functions under appropriate conditions

Limitations

- ▶ **Step size sensitivity:** Performance depends critically on step size selection
- ▶ **Local convergence:** May converge to local minima in non-convex problems
- ▶ **Fixed step size:** No adaptation to problem characteristics


Computational Performance I

Computational Performance II

Algorithm Complexity Visualization

[@fig:complexity] provides a visualization of the algorithm's computational characteristics, including time and space complexity analysis across different problem scales.

Algorithm Performance Analysis
Gradient Descent Convergence Characteristics



Validation I

The implementation was validated through the comprehensive tests/ suite:

- ▶ **Integration tests** verifying algorithm convergence and visualization pipelines.
- ▶ **Infrastructure tests** covering all underlying mechanisms across `infrastructure.reporting`, `infrastructure.validation`, and `infrastructure.rendering`.
- ▶ **Numerical accuracy** checks verified systematically using PyTest.

All tests pass with coverage exceeding the 90% threshold, ensuring implementation correctness across core logic, convergence detection, and logging pathways without the use of mocks.

Discussion I

The experimental results validate the gradient descent implementation and confirm the theoretical convergence predictions from [@sec:methodology]. The monotonic relationship between step size and iteration count ([@tbl:opt_results]) aligns with the convergence factor analysis in [@eq:convergence_factor], while the uniform solution accuracy across all step sizes demonstrates the robustness of the convergence criterion $\|\nabla f(x)\| < \epsilon$. The automated analysis pipeline successfully generated six publication-quality visualizations and structured numerical outputs, validating the template's end-to-end research workflow from algorithmic implementation through automated infrastructure-driven reporting and manuscript integration.

Discussion II

As a meta-architectural note: the perfect embedding of these outputs into this document, including all dynamic references (e.g., [@fig:stability]), confirms the absolute reliability of the `infrastructure/rendering/pdf_renderer.py` module handling the Pandoc conversion.