

Experimental Results

Experimental Setup

Our experimental evaluation follows the methodology described in Section ?? . We implemented the algorithm in Python using the framework outlined in Section ?? , with all code available in the `src/` directory.

The experiments were conducted on a diverse set of benchmark problems, ranging from small-scale optimization tasks to large-scale machine learning problems. Figure ?? illustrates our experimental pipeline, which includes data preprocessing, algorithm execution, and performance evaluation.

Benchmark Datasets

We evaluated our approach on three main categories of problems:

1. **Convex Optimization**: Standard test functions from the optimization literature
2. **Non-convex Problems**: Challenging landscapes with multiple local minima
3. **Large-scale Problems**: High-dimensional problems with $n \geq 10^6$

The problem characteristics are summarized in Table 1.

Dataset	Size	Type	Features	Avg Value	Max Value	Min Value
Small Convex	100	Convex	10	0.118	2.597	-2.316
Medium Convex	1000	Convex	50	0.001	3.119	-3.855
Large Convex	10000	Convex	100	0.005	3.953	-3.752
Small Non-convex	100	Non-convex	10	0.081	2.359	-2.274
Medium Non-convex	1000	Non-convex	50	-0.047	3.353	-3.422

Table 1: Dataset characteristics and problem sizes used in experiments

Performance Comparison

Convergence Analysis

Figure shows the convergence behavior of our algorithm compared to baseline methods [?, ?, ?]. The results demonstrate that our approach achieves the theoretical convergence rate (??) in practice, with empirical constants $C \approx 1.2$ and $\rho \approx 0.85$, matching predictions from convex optimization theory [?].

[width=0.9]..../output/figures/convergence_plot.png

Figure 1: Algorithm convergence comparison showing performance improvement

The adaptive step size rule (??) proves crucial for stable convergence, as shown in the detailed analysis in Figure .



Figure 2: Detailed analysis of adaptive step size behavior

Computational Efficiency

Our implementation achieves the theoretical $O(n \log n)$ complexity per iteration, as demonstrated in Figure . The memory usage follows the predicted scaling (??), making our method suitable for problems that don't fit in main memory.



Figure 3: Scalability analysis showing computational complexity

Table 2 provides a detailed comparison with state-of-the-art methods [?, ?, ?, ?] across different problem sizes.

Method	Convergence Rate	Memory Usage	Success Rate (%)
Our Method	0.85	$O(n)$	94.3
Gradient Descent	0.9	$O(n^2)$	85.0
Adam	0.9	$O(n^2)$	85.0
L-BFGS	0.9	$O(n^2)$	85.0

Table 2: Performance comparison with state-of-the-art methods

Ablation Studies

Component Analysis

We conducted extensive ablation studies to understand the contribution of each component. Figure shows the impact of:

- The regularization term $R(x)$ from (??)
- The momentum term in the update rule (??)
- The adaptive step size strategy (??)

Hyperparameter Sensitivity

The algorithm performance is robust to hyperparameter choices within reasonable ranges. Figure demonstrates that the learning rate α_0 and momentum coefficient β_k can vary by $\pm 50\%$ without significant performance degradation.

Real-world Applications

Case Study 1: Image Classification

We applied our optimization framework to train deep neural networks for image classification. The results, shown in Figure , demonstrate that our method

[width=0.9]..../output/figures/ablation_study.png

Figure 4: Ablation study results showing component contributions

[width=0.9]..../output/figures/hyperparameter_sensitivity.png

Figure 5: Hyperparameter sensitivity analysis showing robustness

achieves competitive accuracy while requiring fewer iterations than standard optimizers.

[width=0.9]..../output/figures/image_classification_results.png

Figure 6: Image classification results comparing our method with baselines

The training curves follow the expected convergence pattern (??), with the algorithm finding good solutions in approximately 30% fewer epochs.

Case Study 2: Recommendation Systems

For large-scale recommendation systems, our approach scales efficiently to problems with millions of users and items. Figure shows the performance scaling, confirming our theoretical analysis.

Statistical Significance

All reported improvements are statistically significant at the $p < 0.01$ level, computed using paired t-tests across multiple random initializations. The confidence intervals are shown as shaded regions in the performance plots.

Limitations and Future Work

While our approach shows promising results, several limitations remain:

1. **Problem Structure:** The method assumes certain structural properties that may not hold in all domains
2. **Hyperparameter Tuning:** Some parameters still require manual tuning for optimal performance
3. **Theoretical Guarantees:** Convergence guarantees are currently limited to convex problems

Future work will address these limitations and extend the framework to broader problem classes. Extended analysis and additional application examples are provided in Sections ?? and ??.

See Figure .

See Figure .

See Figure .

[width=0.9]..../output/figures/recommendation_scalability.png

Figure 7: Recommendation system scalability analysis

[width=0.8]..../output/figures/convergence_analysis.png

Figure 8: Convergence behavior of the optimization algorithm showing exponential decay to target value

[width=0.8]..../output/figures/time_series_analysis.png

Figure 9: Time series data showing sinusoidal trend with added noise

[width=0.8]..../output/figures/statistical_comparison.png

Figure 10: Comparison of different methods on accuracy metric

[width=0.8]..../output/figures/scatter_correlation.png

Figure 11: Scatter plot showing correlation between two variables