Methods: Synthesis, gaps, and coverage scores

Synthesis (`src/synthesis.py`) maps each handbook section to the tuple of evidence items whose themes intersect the section's `theme_ids`. The **coverage score** is computed by `section_coverage_score()`: the sum of weights for that section, capped at 1.0, so many weak items cannot masquerade as certainty.

Sections whose score is **strictly below** the configured gap threshold (default `DEFAULT_GAP_THRESHOLD = 0.35`) are listed in `gaps`. Callers may pass `synthesize(corpus, gap_threshold=…)` when a stricter or looser policy applies; the chosen value is stored on `SynthesisResult.gap_threshold` and echoed in JSON metrics via `build_metrics_report()` (`gap_threshold`, `gap_section_ids`, `gap_count`).

Figure 1 shows section-level scores for the Riverbend fixture. Figure 2 shows how many evidence rows attach to each theme id in the same fixture. Figure 3 repeats the section ordering used in metrics export (sorted ids) and colors bars by membership in