

Results: Language, Topics, and Embeddings

RQ3: Topical and Linguistic Structure I

Text analysis operates over titles, abstracts, and (when available) full text. A TF-IDF representation over a 137-feature vocabulary feeds non-negative matrix factorization, which extracts 5 latent topics cross-cutting the subfield taxonomy. The top vocabulary terms are: atmospheric, uncertainty, spectra, temperature, abundance, analysis, assumptions, opacity, retrievals, sources, stellar, coverage, evaluated, observed, across, composition, models, provide, competing, retrieval.

Table 3. NMF topics extracted from the corpus.

RQ3: Topical and Linguistic Structure I

Topic	Top terms
0	atmospheric, cloud, retrieved, species, changes, structure, models, explore
1	prior, methane, depend, spectral, dioxide, simulated, carbon, constrained
2	contributions, separate, resolution, planetary, high, systematics, modelled, instrument
3	population, reporting, studies, preprocessing, require, harmonized, separates, specific
4	model, intervals, misspecification, noise, quantify, effect, record, dependencies

RQ3: Topical and Linguistic Structure I

The topic labels are computed from the retained corpus and are not hand-assigned. Their interpretation should be checked against the generated topic-term weights and the configured subfield taxonomy; fixture-derived topics describe synthetic text patterns, not the scientific field.

RQ3: Topical and Linguistic Structure I

NMF Topic — Top Terms





Figure 1: Topic-term bar charts for Exoplanet Atmospheres. Each panel shows the top weighted terms for one of 5 NMF topics, with bar length proportional to the topic-term weight in the $\mathbf{H}$ matrix.

Document Embeddings I


Offline deterministic embeddings (TF-IDF followed by truncated SVD) place every document in a shared 50-dimensional vector space. Embedding the same text twice yields identical vectors, so the derived similarity matrix, nearest-neighbour lists, clusters, and two-dimensional projection are all reproducible.

Document Embeddings I



Document Embeddings II

Figure 2: PCA projection of document embeddings for Exoplanet Atmospheres. Each point represents one document projected onto the first two principal components of the TF-IDF/SVD embedding. Colours indicate subfield assignment, showing how the topical geography relates to the keyword taxonomy.



Document Embeddings IV

Figure 3: Hierarchical clustering dendrogram of document embeddings. The tree shows the similarity structure of the corpus: documents that join low in the tree are semantically similar, while high-level splits separate the major topical clusters.

The TF-IDF term heatmap reveals which terms discriminate between subfields: terms with high between-subfield variance (rather than high global mean) are selected for display.

Term Analysis I




Figure 4: Term heatmap for Exoplanet Atmospheres. Each cell shows the mean TF-IDF weight of a term within a subfield. Terms are selected by between-subfield variance to highlight discriminative vocabulary rather than globally frequent terms.

Term Analysis I

- **Definition:** Term analysis is a linguistic method that examines the meaning and usage of individual words or phrases within a specific context.
- **Purpose:** The primary goal of term analysis is to understand the semantic relationships between words and their context, often used in lexicography, linguistics, and translation studies.
- **Scope:** Term analysis typically focuses on the morphological and syntactic structure of words, as well as their semantic and pragmatic functions.
- **Methodology:** The methodology involves a systematic approach to identifying and analyzing the components of a word or phrase, including its root, affixes, and grammatical forms.
- **Applications:** Term analysis is widely applied in various fields, including lexicography, linguistics, translation studies, and computational linguistics.
- **Challenges:** One of the main challenges in term analysis is the ambiguity of words, where a single word can have multiple meanings depending on the context.
- **Tools:** Various tools and software are available to assist in term analysis, including morphological analyzers, syntactic parsers, and semantic networks.
- **Importance:** Term analysis is a fundamental skill for linguists, lexicographers, and translators, as it provides a deep understanding of the structure and meaning of language.
- **Conclusion:** Term analysis is a complex but essential linguistic tool that helps in understanding the intricate relationships between words and their context.

Named Entity Analysis I

Named entity extraction over the 80 abstracts identified 1 unique entities. The most frequent entities reflect the retained corpus and the configured extraction rules; source coverage and fixture status determine how they may be interpreted.

Table 4. Top named entities in abstracts.

Named Entity Analysis I

Entity	Frequency
--------	-----------

JWST	17
------	----

Named Entity Analysis II

Entity Frequency

JWST 17

Top 1 Named Entities in Abstracts




Figure 5: Top named entities for Exoplanet Atmospheres. The horizontal bar chart shows the 20 most frequently

Named Entity Analysis I

Named Entity Analysis II

Table 5. Top keyphrases by TF-IDF score.

Named Entity Analysis I

Keyphrase	Score
jwst	0.0714
population	0.0556
analyses	0.0435
calibration	0.0435
calibration choices	0.0435
choices	0.0435
combined	0.0435
jwst analyses	0.0435
models	0.0435
models calibration	0.0435
models calibration choices	0.0435
model	0.0400
atmospheric	0.0392
abundance	0.0385
assumptions	0.0385


The TF-IDF/SVD embeddings place every document in a 50-dimensional vector space. K-means clustering with $k = 5$ clusters partitions the corpus into topically coherent groups. The top similar document pairs, ranked by cosine similarity, reveal the most closely related works in the corpus.

Table 6. Top 10 most similar document pairs.

Embedding Similarity and Clustering II

Paper A	Paper B	Similarity
doi:10.5555/exoplanet atmosphe	doi:10.5555/exoplanet atmosphe	1.0000
doi:10.5555/exoplanet atmosphe	doi:10.5555/exoplanet atmosphe	1.0000
doi:10.5555/exoplanet atmosphe	doi:10.5555/exoplanet atmosphe	1.0000
doi:10.5555/exoplanet atmosphe	doi:10.5555/exoplanet atmosphe	1.0000
doi:10.5555/exoplanet atmosphe	doi:10.5555/exoplanet atmosphe	1.0000
doi:10.5555/exoplanet atmosphe	doi:10.5555/exoplanet atmosphe	1.0000
doi:10.5555/exoplanet atmosphe	doi:10.5555/exoplanet atmosphe	1.0000
doi:10.5555/exoplanet atmosphe	doi:10.5555/exoplanet atmosphe	1.0000
doi:10.5555/exoplanet atmosphe	doi:10.5555/exoplanet atmosphe	1.0000
doi:10.5555/exoplanet atmosphe	doi:10.5555/exoplanet atmosphe	1.0000

Top 15 Most Similar Document Pairs



Embedding Similarity and Clustering I

These embeddings support semantic retrieval over the corpus and the visual map of the literature's topical geography.