

## Results and Discussion

mypy-as-oracle proof-of-detection

tests/test\_mypy\_oracle.py runs `mypy --strict` as a real subprocess against every fixture under `tests/mypy_fixtures/` — **six** known-bad programs, one per ISC-2/4/15/18/31/32 — and asserts a **non-zero** exit code plus the presence of a genuine error: line on each (ISC-19, reusing a consumed protocol-phase instance, deliberately has **no** fixture here: it is a runtime-only discipline, not a static one — see sec. 3 — so there is nothing for a static oracle to catch). **Three** known-good fixtures are the suite's positive controls, each guarding a generic-API or structural-conformance surface `src/` itself never instantiates concretely: `good_agent_belief_instantiation.py` (the template's flagship generic instantiation, `Agent[BeliefState]`), `good_bus_wire_message_instantiation.py`

(InProcessBus[WireMessage]), and good\_pheromone\_conformance.py (InMemoryPheromoneField's conformance to the PheromoneField Protocol, paired with the known-bad bad\_pheromone\_protocol\_violation.py above). The same test then runs `mypy --strict` against the real `src/template_formal/tree` and asserts a **zero** exit code (ISC-39): the negative controls prove the oracle can detect the class of bug it is supposed to detect, and the positive controls (the clean `src/` tree, plus the three good-fixture instantiations) prove the shipped code does not contain that class of bug **and** that its generic/structural API surfaces are actually usable the way the manuscript advertises — a distinction this template did not always get right (see below). No fixture in this suite is skipped, xfailed, or hidden behind a suppressing

# type: ignore (ISC-40 anti) — a fixture that stopped failing under mypy --strict would itself be a signal that the oracle had gone vacuous, not that the type architecture had improved.

A cross-vendor audit of this template caught a genuine defect the src/-only oracle could not see: an earlier revision of the GaussianBelief structural Protocol (sec. 3) declared its mean/variance members as plain mutable attributes, which a frozen=True dataclass can never satisfy under mypy --strict — so Agent[BeliefState], the template's own reference instantiation, failed mypy --strict at the one call site that matters, even while the src/-only gate reported zero errors (because src/ never itself instantiates Agent with a concrete type argument). The fix was to declare GaussianBelief.mean/.variance as read-only

@property members instead of plain attributes, and the good-fixture above is the permanent regression guard. This episode is recorded here rather than quietly fixed and forgotten because it is precisely the failure mode sec. 3 warns about in the abstract: a green gate over an incomplete scan-set is not the same claim as “the type architecture works,” and this template’s own build process needed a second, adversarial pass to catch the gap between the two.

Protocol fault-injection negative controls

## Protocol fault-injection negative controls I

`tests/network/test_bus.py` and `tests/network/test_handshake_over_bus.py` drive a real handshake through the in-process bus (`network/bus.py`) with each fault mode enabled independently. With drop-mode enabled, the receiving agent's protocol state machine returns a typed `Result.Err(ProtocolViolation(...))` rather than crashing or silently advancing phase (ISC-23). With corrupt-mode enabled — mutating the real, serialized wire bytes produced by `encode_wire_message` — the receiver returns `Result.Err(MalformedMessage(...))` (ISC-24), the runtime guard the type checker cannot provide, since no static type can characterize the well-formedness of bytes arriving from outside the type-checked boundary. A deterministic-seed test (`test_fault_injector_determinism_same_seed_same_sequence`) confirms the same seed reproduces the identical fault sequence (ISC-22),

## Protocol fault-injection negative controls I

which is what makes these fault-injected runs reproducible evidence rather than occasionally-flaky noise that happens to pass. Every phase-transition claim in sec. 3 has a paired fault-injected test alongside its happy-path test (ISC-25 anti) — there is no protocol test in this suite that only exercises the fault-free path.

Colony convergence: a real stigmergic mechanism, honestly scoped

## Colony convergence: a real stigmergic mechanism, honestly scoped I


`tests/colony/test_colony_integration.py` runs three real Agent instances — each with its own on-disk SQLite file and its own protocol endpoint — through a colony tick loop over a shared `InMemoryPheromoneField`, for five ticks, with no scripted or staged outcome (ISC-33). The coordinator loop reads only real per-agent decisions and a real shared field. Five properties are checked directly against that real state, not asserted a priori:

## Colony convergence: a real stigmergic mechanism, honestly scoped I

1. All three agents resolve the same free-energy tie to the same first location on tick 0. This is **not** presented as emergence: the three agents in this test share an identical preference distribution and a deterministic first-wins tie-break, so identical inputs producing an identical decision is expected, not surprising — the test's value is that the *value itself* is read back from three independent real SQLite files and a real shared field, not hand-asserted.
2. That agreement persists every subsequent tick.
3. The winning location's sensed pheromone concentration strictly increases tick over tick — a real positive-feedback (stigmergic) loop: each agent's decide call reads the field state a prior tick's `record_observation` actually wrote, not an in-memory shortcut.
4. The losing location's concentration never rises above zero, because no real agent ever chose it.
5. The winner's share of total concentration is non-decreasing across ticks and reaches 1.0 by the final tick — computed from real per-agent state.

## Colony cc

Deterministic demo colony (3 agents, 5 ticks)



## Colony convergence: a real stigmergic mechanism, honestly scoped I

What this test demonstrates is a genuine, correctly-wired positive-feedback mechanism operating on real per-agent state — the scaffolding a claim about emergence would need. It does **not** demonstrate emergence in the stronger sense (a colony-level regularity that is *not* a direct consequence of individual agent design): with three identical, deterministic agents and a symmetric tie-break, the tick-0 agreement is guaranteed by construction, not a discovered regularity. A template that wanted to earn the word “emergent” would need heterogeneous agent preferences or a stochastic tie-break and would need to show consensus still forms despite that heterogeneity — that is future work relative to *this specific test*, but it is exactly the claim the next subsection earns with a real  $N=150$  statistical test, not merely asserted here.

Colony convergence under heterogeneity: a genuinely earned statistical result

## Colony convergence under heterogeneity: a genuinely earned statistical result I

`tests/colony/test_colony_convergence_statistics.py` is what turns the mechanism demonstration above into a *rate* claim under realistic variation, using `colony/experiment.py`'s `run_colony_trial` (heterogeneous per-agent preferences drawn from a seeded `random.Random`, plus seeded Gaussian sensing noise — see sec. 3's module layout) and `colony/stats.py`'s closed-form `wils_on_score_interval`:

## Colony convergence under heterogeneity: a genuinely earned statistical result I

### Hypothesis (Calibrated-configuration convergence floor)

*The true colony convergence rate, at the calibrated baseline configuration (8 heterogeneous agents, 2 locations, 30 ticks, sensing-noise  $\sigma = 0.5$ , decay = 0.46), is at most 0.8. Rejected at  $\alpha = 0.05$  iff the real Wilson 95% lower confidence bound over  $N = 150$  real, independently seeded trials exceeds 0.8.*

The actual run settles this hypothesis: **140/150 trials converged (rate = 0.9333), Wilson 95% CI = (0.8816, 0.9634)** — the lower bound clears 0.8 by a wide margin, not by a hair.

## Colony convergence under heterogeneity: a genuinely earned statistical result I

**Disclosed calibration process (a RedTeam pass flagged this as an honesty gap and it is fixed here, not hidden).** tests/colony/test\_colony\_convergence\_statistics.py's own docstring states plainly that `decay=0.46/sensing_noise_std=0.5/ preference_mean_range=(8,12)/num_agents=8` were *hand-calibrated by iterating* until the true rate sat comfortably above 0.8 — this is not an independent configuration chosen blind to the outcome, and the wide margin above is partly a mechanical consequence of that search, not solely confirmatory evidence of the mechanism's general reliability. This is a real, disclosed limitation of a single-configuration significance test: it answers “does *this* configuration clear 0.8” honestly, but it does **not**, by itself, answer “how does the convergence rate behave as configuration

## Colony convergence under heterogeneity: a genuinely earned statistical result I

parameters vary” — which is precisely the gap `sec. sec:results-discussion`’s “Eight pre-registered analyses” subsection below closes with a real parameter sweep across six independent decay values, run and reported without retuning any of them to hit a target. Two guards keep the single-point claim honest rather than vacuous: a negative control re-running the *original* fully-symmetric, zero-noise configuration through this same harness reproduces exact 100% convergence (the “guaranteed by construction” claim above is a real, reproducible fact under the harness that also produces the  $> 0.8$  claim, not a coincidence of two unrelated code paths), and a positive-control-that-can-fail — near-total pheromone evaporation (decay = 0.97) plus sensing noise ( $\sigma = 4.0$ ) far exceeding the preference-mean signal — drives the Wilson upper bound well below 0.5, proving the  $> 0.8$

gate is not vacuously satisfiable for every configuration.

**Confound, honestly disclosed:** this positive control changes decay *and* sensing\_noise\_std simultaneously, so on its own it cannot attribute the collapse to loss of the stigmergic mechanism specifically versus noise magnitude alone. The decay-only sweep below (holding noise fixed at  $\sigma = 0.5$ ) is the experiment that isolates decay as a single variable and reports what it alone does to convergence; the noise-magnitude axis in isolation (varying sensing\_noise\_std alone at a fixed decay) remains untested and is still named explicitly as future work, not silently generalized past. A genuine deposit\_amount=0 ablation *is* run elsewhere in this section — see Experiment B's zero-deposit condition below — but at the calibrated baseline configuration (decay=0.46, sensing\_noise\_std=0.5), not at this

## Colony convergence under heterogeneity: a genuinely earned statistical result I

extreme positive-control configuration; it closes the attribution gap for Experiment B's real-vs-null comparison specifically, not for this positive control's decay/noise conflation.

**Half of this confound, now closed: does pheromone still matter once decay and noise are already this severe?** A dedicated ablation (Experiment I, `tests/colony/test_colony_experiments_extended.py`) runs the missing `deposit_amount=0` condition *at this exact extreme configuration* (decay=0.97, sensing\_noise\_std=4.0, identical n=50/seed\_base=0 as the positive control above), asking a narrower, precisely-stated question:

## Hypothesis (Pheromone contribution under severe decay and noise)

*The real mechanism with pheromone deposit disabled ( $deposit\_amount=0.0$ ) at the extreme positive-control configuration ( $decay=0.97$ ,  $\sigma_{sense}=4.0$ ,  $n=50$ ,  $seed\_base=0$ ) converges no differently than the existing full-mechanism ( $deposit\_amount=1.0$ ) positive control at the identical configuration. Falsified by a two-sided Fisher's exact test on the pairwise comparison reaching significance at  $\alpha=0.05$ , together with the full mechanism's Wilson lower bound exceeding the zero-deposit condition's Wilson upper bound.*

Colony convergence under heterogeneity: a genuinely earned statistical result I

**Real result:**

## Colony convergence under heterogeneity: a genuinely earned statistical result I

| condition                           | successes/n | rate   | Wilson 95% CI    |
|-------------------------------------|-------------|--------|------------------|
| full mechanism (deposit_amount=1.0) | 9/50        | 0.1800 | (0.0977, 0.3080) |
| zero-deposit (deposit_amount=0.0)   | 0/50        | 0.0000 | (0.0000, 0.0713) |

## Colony convergence under heterogeneity: a genuinely earned statistical result I

A two-sided Fisher's exact test on this pairwise comparison (the correct small-sample test here, matching this section's own precedent at the decay sweep's 100%-boundary dip, since one group sits exactly at 0%) gives  $p = 0.0026$  — significant at  $\alpha = 0.05$ . The full mechanism's Wilson lower bound (0.0977) also clears the zero-deposit condition's Wilson upper bound (0.0713), though narrowly.  $H_0$  is **rejected**: even at this deliberately-defeated configuration, the pheromone/stigmergic channel still contributes something real and statistically distinguishable from having no pheromone channel at all. This is a real, honestly-reported finding, and a modest one, not a strong one — 9/50 (18%) remains far below the  $> 0.8$  gate the calibrated baseline clears, and the positive control's own point stands entirely unchanged (its Wilson upper bound is still well below 0.5; the mechanism is still

## Colony convergence under heterogeneity: a genuinely earned statistical result I

severely, deliberately defeated by this configuration). **What this does and does not establish:** it establishes that the near-total collapse the positive control demonstrates is not the *whole* story of “noise/decay alone, pheromone irrelevant” — some residual stigmergic signal survives even under decay this close to total evaporation and noise this far past the preference-mean signal. It does **not**, by itself, decompose how much of the original collapse is attributable to decay=0.97 alone versus sensing\_noise\_std=4.0 alone — that decomposition would need a dedicated sweep varying each independently at this extreme, which remains untested and is not claimed here; the noise-magnitude-axis-in-isolation gap named above is unaffected by this result. Gated by test\_extreme\_zero\_deposit\_pheromone\_channel\_still\_contributes\_under\_severe\_decay\_and\_noise

Colony convergence under heterogeneity: a genuinely earned statistical result I

in `test_colony_experiments_extended.py`.

## Colony convergence under heterogeneity: a genuinely earned statistical result I




Figure 2: Convergence-tick distribution from this template's own demo statistics sweep (`src/template_formal/colony/demo.py::run_statistics_sweep`,  $N = 40$  trials, same real, heterogeneous/noisy configuration as the  $N = 150$  test above, just smaller so the demo script stays fast): 37/40 trials converged (92.5%). Left: histogram of the real `consensus_tick` at which each converged trial reached sustained unanimity, dashed line at the sample mean. Right: the same distribution as an empirical CDF. This demo figure is illustrative of the *shape* of the real distribution at a smaller, faster  $N$  — the  $N = 150$  Wilson-interval numbers quoted above, not this  $N = 40$  run, are what the paper's statistical claim is bound to.

## Colony convergence under heterogeneity: a genuinely earned statistical result I

A companion test (`test_colony_coordinator_never_touches_an_agents_internal_storage_or_protocol_handle`) asserts the coordinator's only reachable surface is each Agent's eight public members and the PheromoneField Protocol's three public methods — making ISC-34's “the coordinator cannot reach an agent's internals” anti-criterion a structural fact about the object graph, not a claim about what the test *happened* to do.

## Colony convergence under heterogeneity: a genuinely earned statistical result I

The same  $N=150$  batch also feeds `colony/stats.py`'s `pearson_r` in one further, deliberately weak, exploratory check: the correlation between each converged trial's preference-mean spread and its `consensus_tick` is  $r = 0.1648$  — asserted only as finite and not implausibly close to  $\pm 1$  ( $-0.99 < r < 0.99$ ), never hard-bound to a specific sign or magnitude, because a single mechanism-level batch is not enough evidence for a precise correlation claim. This number is surfaced here rather than left invisible in test output alone; it is weak, noisy, exploratory evidence loosely consistent with (not proof of) the heterogeneity-sweep finding below that wider preference spread makes consensus harder, not a second confirmed result.

Eleven pre-registered analyses across three experiment families

## Eleven pre-registered analyses across three experiment families I

The  $N=150$  statistical claim above answers one question — does the real mechanism converge reliably at *one* calibrated configuration — and deliberately leaves three experiment families open: whether that rate is sensitive to the pheromone-evaporation rate, whether the mechanism actually beats a random-choice colony with no mechanism at all, and whether heterogeneity's effect is monotonic or something stranger. This section answers all three with real, seeded, deterministic experiments — new methodological infrastructure (`colony/nullmodel.py`, a random-choice baseline harness, and `colony/sweep.py`, a generic parameter-sweep runner reusing `run_colony_trial` and the same Wilson-interval machinery) rather than one-off scripts — reported in `tests/colony/test_colony_experiments_extended.py`, whose assertions are pinned to the exact successes counts quoted

Eleven pre-registered analyses across three experiment families I

below (every run is fully deterministic given its seed sequence, so these are not approximate).

## Experiment A — does convergence rate vary monotonically with decay? I

### Hypothesis (Monotonic decay sensitivity)

*The colony's convergence rate is a monotonically non-decreasing function of decay over  $\{0.10, 0.30, 0.46, 0.60, 0.80, 1.00\}$ , holding every other parameter at the calibrated baseline (8 agents, 2 locations, 30 ticks,  $\sigma_{sense} = 0.5$ , preference range (8, 12)).*

*Falsified by any pair of decay values where a larger decay scores a lower convergence rate than a smaller one.*

Experiment A — does convergence rate vary monotonically with decay? I

**Real result** (n=60 independently-seeded trials per value, seed\_base=0):

## Experiment A — does convergence rate vary monotonically with decay? I

| decay | successes/n | rate   | Wilson 95% CI    |
|-------|-------------|--------|------------------|
| 0.10  | 0/60        | 0.0000 | (0.0000, 0.0602) |
| 0.30  | 0/60        | 0.0000 | (0.0000, 0.0602) |
| 0.46  | 56/60       | 0.9333 | (0.8407, 0.9738) |
| 0.60  | 60/60       | 1.0000 | (0.9398, 1.0000) |
| 0.80  | 60/60       | 1.0000 | (0.9398, 1.0000) |
| 1.00  | 56/60       | 0.9333 | (0.8407, 0.9738) |

## Experiment A — does convergence rate vary monotonically with decay? I

$H_0$  is **rejected**: decay  $\in \{0.60, 0.80\}$  both reach 100% while decay = 1.00 (total pheromone evaporation every tick) drops back to 56/60 — the two counts at 0.46 and 1.00 are identical, 56/60, so this is not simply “more decay is always better.” The shape is not a smooth trend at all: it is a **threshold**, not a monotonic slope. Below decay  $\approx 0.35$  the colony essentially never converges (both 0.10 and 0.30 score exactly 0/60, Wilson upper bound under 0.06); a finer, exploratory sweep at the same seed base (13 points, 0.10 through 1.00, not gated by a regression test) shows the transition is sharp — 0% at decay  $\leq 0.30$ , rising through 18% (0.38), 57% (0.42), 93% (0.46) to reach the 100% plateau by decay  $\approx 0.60$ .

## Experiment A — does convergence rate vary monotonically with decay? I

**Precision correction (a second Forge cross-vendor pass on this section found the wording above originally overreached, and it is corrected here):** the entire non-monotonicity finding formally rests on one pairwise gap — 60/60 at decay  $\{0.60, 0.80\}$  versus 56/60 at decay = 1.00 — whose Wilson intervals, (0.9398, 1.0000) and (0.8407, 0.9738), overlap heavily. A **Fisher's exact test** on this one pair (the correct small-sample test here, not a normal-approximation two-proportion  $z$ , precisely because one group sits exactly at the 100% boundary — the same boundary pathology `wilson_score_interval` itself was hardened against, ISC-82) gives two-sided  $p = 0.1187$  (`colony/stats.py`'s new `fisher_exact_test_two_sided`, a real hypergeometric computation via `math.comb`, not an approximation): at this single seed base, this one pairwise comparison alone does **not**

## Experiment A — does convergence rate vary monotonically with decay? I

clear conventional significance. The honest reason the decline is still reported as real, not as this template overreaching its own evidence, is that no single pairwise test carries the claim: the 13-point exploratory sweep's monotonically *rising* shape below the plateau, together with the cross-seed replication of the dip at seed bases 1000 and 5000 (sec. sec:results-discussion, “Honesty hedges” below), is where the actual evidentiary weight sits.

## Experiment A — does convergence rate vary monotonically with decay? I

**Ordered-trend test (Cochran–Armitage), now added.** A prior Forge cross-vendor pass named a formal trend test across the *full* ordered sweep (rather than one boundary-adjacent pairwise comparison) as the correctly-scoped next step; this round adds it. The **Cochran–Armitage test for trend** (`colony/stats.py`'s new `cochran_armitage_trend_test`, a `stdlib`-only closed-form  $Z$ -statistic from `math` and `statistics.NormalDist` — the textbook ordered-groups test for a *binary* endpoint, no `scipy`) over the full ordered decay sweep — successes  $\{0, 0, 56, 60, 60, 56\}$  out of 60 at scores  $\{0.10, 0.30, 0.46, 0.60, 0.80, 1.00\}$  — gives  $Z = +14.57$ , two-sided  $p < 10^{-10}$ : a strongly significant *overall increasing* association between decay and convergence. **This does not overturn the local non-monotonicity, and the two results are deliberately reported side by side.** Cochran–Armitage and Fisher's

## Experiment A — does convergence rate vary monotonically with decay? I

exact answer different questions: the trend test detects the broad low-decay-floor-to-high-decay-plateau *rise* (which dominates the linear signal and swamps the small interior dip), while the pairwise Fisher test above isolates the *local* 60/60-versus-56/60 dip at the very top of the range ( $p = 0.1187$ , not conventionally significant). A significant positive trend statistic and a real, separately-established interior dip coexist — the sweep rises overall *and* dips at its extreme — so the larger effect is not permitted to silently erase the smaller one. Both facts are gated in `test_colony_experiments_extended.py` (`test_cochran_armitage_finds_a_strong_positive_trend_across_the_full_decay_sweep`, plus a companion test asserting the raw dip still holds alongside the positive trend).

## Experiment A — does convergence rate vary monotonically with decay? I

### **A plausible mechanism, first from trace inspection alone.**

Manual inspection of individual trial traces at low decay (e.g. decay = 0.10) suggested why: with too little evaporation, both locations' sensed concentrations grow together, roughly in lockstep, tick over tick, past 30 within thirty ticks — far outside the agents' preference range of (8, 12). Once both candidates are similarly far from every agent's preference, the free-energy KL term's discriminating signal between them shrinks to the same order of magnitude as the sensing noise ( $\sigma = 0.5$ ), and agents effectively split unpredictably between locations tick to tick rather than reinforcing one attractor — visible directly in the trace as persistent 50/50-ish oscillation rather than settling. When this was first written, this was explicitly flagged as a plausible causal account consistent with what the traces show, **not** a claim independently isolated by a

## Experiment A — does convergence rate vary monotonically with decay? I

dedicated ablation (e.g. artificially capping sensed concentration) — that was named as a distinct future experiment. It has now been built.

### Hypothesis (Sensed-concentration cap recovers low-decay convergence)

*Capping sensed concentration does not improve convergence at low decay ( $\text{decay} = 0.10$ ,  $\text{decay} = 0.30$ ) relative to the uncapped baseline (0/60 at both). Falsified by the capped condition's Wilson lower bound clearing the uncapped condition's Wilson upper bound at both decay values.*

## Experiment A — does convergence rate vary monotonically with decay? II

ColonyTrialConfig gained a new, opt-in  
sensed\_concentration\_cap: float | None = None field  
(None is fully behavior-preserving — every existing test and result  
in this paper, none of which sets it, reproduces byte-for-byte  
identically), wired into run\_colony\_trial's per-tick candidate  
loop so that when set, the sensed concentration field.sense(lo  
cation) is clipped to at most the cap *before* the sensing-noise  
term is added — a saturating-sensor

## Experiment A — does convergence rate vary monotonically with decay? I

model, not reduced measurement noise. The cap value, 13.0, is chosen just above the preference range's ceiling (12) — only 1.0 above it, two sensing-noise standard deviations ( $\sigma = 0.5$ ) — high enough that ordinary within-range sensing early in a trial is unaffected, low enough to bound exactly the runaway past-preference-range growth the hypothesis describes. This was not selected by sweeping many cap values and reporting only the best one: an exploratory, ungated probe found the effect holds essentially unchanged for any cap in  $[12.5, 14.0]$  and fades out entirely by  $\text{cap} = 20.0$  (which never binds within the 30-tick horizon and reproduces the uncapped 0/60 exactly) — 13.0 is a representative point inside the range where the cap has a real, non-degenerate effect, not a cherry-picked edge.

Experiment A — does convergence rate vary monotonically with decay? I

**Real result** (n=60 independently-seeded trials per value, identical seed\_base=0, identical calibrated baseline configuration except seed\_concentration\_cap=13.0):

## Experiment A — does convergence rate vary monotonically with decay? I

| decay | condition     | successes/n | rate   | Wilson 95% CI    |
|-------|---------------|-------------|--------|------------------|
| 0.10  | uncapped      | 0/60        | 0.0000 | (0.0000, 0.0602) |
| 0.10  | capped (13.0) | 60/60       | 1.0000 | (0.9398, 1.0000) |
| 0.30  | uncapped      | 0/60        | 0.0000 | (0.0000, 0.0602) |
| 0.30  | capped (13.0) | 60/60       | 1.0000 | (0.9398, 1.0000) |

## Experiment A — does convergence rate vary monotonically with decay? I

$H_0$  is **rejected**, decisively: capping alone flips both decay values from near-certain non-convergence to certain convergence — the uncapped Wilson upper bound (0.0602) does not come close to overlapping the capped Wilson lower bound (0.9398) at either decay value. This is real evidence, not merely trace-inspection inference, that the hypothesized mechanism is at least *sufficient* to explain the low-decay failure: removing only the one effect the mechanism identifies (unbounded sensed-concentration growth past the preference range), while changing nothing else about decay, sensing noise, preferences, or the decision rule, is enough on its own to recover convergence. Reported honestly regardless of which way this went (per this manuscript's own discipline): had capping left the rate near zero, that would have refuted or left unconfirmed the mechanistic account as

## Experiment A — does convergence rate vary monotonically with decay? I

stated, and this section would say so plainly rather than force a positive spin — it did not happen here.

**What this does and does not establish.** It establishes sufficiency, at this one configuration and this one cap value, of removing unbounded sensed-concentration growth to fix low-decay non-convergence. It does *not* establish necessity — this ablation does not rule out that some other, untested intervention could also recover convergence at low decay via a different causal path. A real, gated dose-response sweep across cap values (below) now exists — replacing the informal, ungated probe this paragraph previously deferred to — but it covers only  $\text{decay}=0.10$ , not the full decay range, so it does not characterize how the dose-response curve itself shifts as decay varies. This ablation also does not by itself prove the exact step-by-step causal story in the first paragraph above (KL-term signal

## Experiment A — does convergence rate vary monotonically with decay? I

shrinking to the sensing-noise order of magnitude) — it shows that the gross behavioral prediction of that story (bounding growth restores convergence) holds, which is strong corroborating evidence for the account, but the intermediate mechanistic detail (the specific KL-term argument) is still inferred from traces, not measured directly by a dedicated information-theoretic instrument. Gated by `test_capping_sensed_concentration_recovers_convergence_at_low_decay` in `test_colony_experiments_extended.py`.

## Experiment A — does convergence rate vary monotonically with decay? I

The manipulation of decay itself, meanwhile, is a real controlled variable in a real simulation (every other input, including the per-agent preference draws and the sensing-noise sequence, is held identical across the swept values via a paired-seed design in `colony/sweep.py`), so the *effect* claim above (decay changes the rate, non-monotonically) was always on firmer ground than the *mechanism* claim (why it changes) — and the mechanism claim is now itself corroborated by a dedicated ablation rather than resting on trace inspection alone.

## Experiment A — does convergence rate vary monotonically with decay? I

**Gating the dose-response curve itself.** The single cap value (13.0) above answers “does capping help at all?” but not “how does the effect change as the cap is relaxed?” — the previous round’s justification paragraph named an *informal, ungated probe* for that question (“the effect holds essentially unchanged for any cap in [12.5, 14.0] and fades out entirely by cap = 20.0 ... not gated, not quoted as a result”). This round promotes that probe into a real, pre-registered, gated sweep.

## Experiment A — does convergence rate vary monotonically with decay? I

### Hypothesis (Sensed-concentration-cap dose-response at low decay)

*Convergence rate at decay = 0.10 does not vary as sensed\_concentration\_cap increases from near the preference-range ceiling (12.5) toward a value that never binds within the tick horizon (20.0). Falsified by any non-monotonic reversal in the ordered sweep, or by no measurable variation at all.*

colony/sweep.py's run\_parameter\_sweep is generic over any ColonyTrialConfig field except seed — confirmed directly rather than assumed, so sensed\_concentration\_cap required zero new sweep machinery, only a direct call. **Real result** ( $n = 60$  per value, identical seed\_base=0, decay=0.10 fixed, otherwise the identical calibrated baseline):

## Experiment A — does convergence rate vary monotonically with decay? I

| cap  | successes/n | rate   | Wilson 95% CI    |
|------|-------------|--------|------------------|
| 12.5 | 60/60       | 1.0000 | (0.9398, 1.0000) |
| 13.0 | 60/60       | 1.0000 | (0.9398, 1.0000) |
| 15.0 | 41/60       | 0.6833 | (0.5577, 0.7869) |
| 16.0 | 13/60       | 0.2167 | (0.1312, 0.3362) |
| 17.0 | 3/60        | 0.0500 | (0.0171, 0.1370) |
| 18.0 | 1/60        | 0.0167 | (0.0029, 0.0886) |
| 20.0 | 0/60        | 0.0000 | (0.0000, 0.0602) |

## Experiment A — does convergence rate vary monotonically with decay? I

$H_0$  (no variation) is **rejected**: the rate declines monotonically (non-increasing at every step — a Cochran–Armitage trend test over the full sweep gives  $Z = -16.4245$ ,  $p < 10^{-10}$ ; a two-sided Fisher's exact test on the plateau (13.0) versus the floor (20.0) gives  $p \approx 2.07 \times 10^{-35}$ , the closed-form value for this perfectly separated 60/60-vs-0/60 table,  $2/\binom{120}{60}$ ). But **the real shape is not the smooth, gradual fade the informal probe's phrasing might suggest**. It is a plateau (100% at  $\text{cap} \in \{12.5, 13.0\}$ ) followed by a steep decline that is essentially *complete by cap* = 18.0 (98%+ of the total plateau-to- floor decline) — a full 2.0 cap units *before* the informally-named  $\text{cap} = 20.0$  “fade out” point, which merely reconfirms a floor already reached two units earlier and reproduces the independently-computed, uncapped  $\text{decay}=0.10$  baseline's 0/60

## Experiment A — does convergence rate vary monotonically with decay? I

exactly. The transition is front-loaded and largely complete well inside the previously-probed range, not a late collapse arriving only near the named endpoint. Cross-checked against Experiment A's own already-gated single point: this sweep's cap = 13.0 point (swept over sensed\_concentration\_cap at fixed decay=0.10) and the single-point ablation's decay=0.10 point (swept over decay at fixed sensed\_concentration\_cap=13.0) describe the identical underlying configuration and reproduce the identical 60/60 exactly, not merely approximately. Gated by eight tests in `test_colony_experiments_extended.py` (`test_cap_dose_response_*`).

## Experiment A — does convergence rate vary monotonically with decay? I

**A cross-vendor audit caught a real bug in `fisher_exact_test_two_sided` itself, exercised for the first time by this experiment's plateau-vs-floor comparison, and it is fixed here, not hidden.** The function's two-sided tail used an additive tolerance (`observed_p + 1e-10`), safe only when `observed_p` is itself around  $1e-10$  or larger. The plateau (13.0) versus floor (20.0) table here is perfectly separated (60/60 vs 0/60), whose true hypergeometric `observed_p` is  $\approx 10^{-35}$  — the additive  $1e-10$  term completely swallowed it, degenerating the two-sided sum to “every table with probability at most  $10^{-10}$ ” instead of “every table at least as extreme as observed,” inflating the reported p-value by 24 orders of magnitude (from the true  $2.07 \times 10^{-35}$  to a wrong  $4.31 \times 10^{-11}$ ). The scientific conclusion is

## Experiment A — does convergence rate vary monotonically with decay? I

unchanged under both values — the table is overwhelmingly significant either way — but a function this manuscript explicitly advertises as “a real hypergeometric computation... not an approximation” owes exactness at every magnitude, not just the ones its own five previously-pinned test cases happened to probe (all of which had observed\_p comfortably above  $10^{-10}$ , so none exercised this failure mode). Fixed to a relative tolerance (`observed_p * (1.0 + 1e-9)`), independently re-verified against `scipy.stats.fisher_exact` and the closed-form  $2/\binom{120}{60}$  for this exact table, and pinned by a new dedicated regression test (`test_fisher_exact_perfectly_separated_table_matches_hand_computed_value` in `tests/colony/test_colony_stats_unit.py`) specifically for a highly-separated table, so a regression back to an additive tolerance

Experiment A — does convergence rate vary monotonically with decay? I

fails immediately rather than only surfacing as a silently wrong number several call-sites away in a manuscript quote — which is exactly how this one survived two prior audit rounds undetected.

## Experiment B — does the real mechanism beat a random-choice baseline? I

### Hypothesis (Mechanism beats random-choice baseline)

*The real stigmergic mechanism's convergence rate is no higher than a colony of agents that ignore the pheromone field entirely and each pick a location uniformly at random every tick (colony/nullmodel.py's run\_null\_model\_trial — no Agent, no BeliefState, no free-energy computation, no pheromone field: literally random.Random(seed).choice(locations) per agent per tick, proven structurally isolated from the rest of the colony machinery by a source-text grep test). Falsified by the real mechanism's Wilson lower bound failing to exceed the null model's Wilson upper bound.*

Experiment B — does the real mechanism beat a random-choice baseline? I

**Real result** ( $N = 150$  trials each, identical seed sequence  
seed\_base=0, calibrated baseline configuration, decay=0.46):

## Experiment B — does the real mechanism beat a random-choice baseline? I

| harness                    | successes/N rate | Wilson 95% CI           |
|----------------------------|------------------|-------------------------|
| real mechanism             | 140/150          | 0.9333 (0.8816, 0.9634) |
| null model (random choice) | 1/150            | 0.0067 (0.0012, 0.0368) |

## Experiment B — does the real mechanism beat a random-choice baseline? I

$H_0$  is **decisively rejected**: the two Wilson intervals do not overlap — the real mechanism's lower bound (0.8816) clears the null model's upper bound (0.0368) by a wide margin. This is the comparison the mechanism-demonstration test and the  $N=150$  statistical test above never made: neither, on its own, rules out “maybe 93% convergence is just what happens with 8 agents and 2 locations regardless of mechanism.” It is not. The null model's own rate is itself a sanity check on `find_sustained_consensus_tick`'s definition: with 8 independent random choosers over 2 locations for 30 ticks, a back-of-envelope expected count of sustained-consensus events is  $150 \times 2 \times (1/2)^8 \approx 1.2$  — consistent with the observed  $1/150$ , i.e. the null model is behaving exactly as an uninformed random baseline should, not as an accidentally-biased one.

## Experiment B — does the real mechanism beat a random-choice baseline? I

**Confound, honestly disclosed and now closed.** A RedTeam pass (round 4) flagged that the comparison above cannot, by itself, attribute the real mechanism's 93%-vs-0.7% gap to the pheromone/stigmergic feedback channel specifically, as opposed to some general noise-robustness of the free-energy decision rule that would exist even without a pheromone field. The null model is not a controlled ablation of the real mechanism — it has no Agent, no BeliefState, no free-energy computation at all; it is a structurally different code path (`random.Random(seed).choice(locations)` per tick), so a real-vs-null comparison alone cannot isolate *which part* of the real mechanism drives the advantage. Closing this required building the missing middle condition: the real Agent/BeliefState/free-energy decision loop, run through the identical `run_colony_trial` harness, but with

## Experiment B — does the real mechanism beat a random-choice baseline? I

`deposit_amount=0.0` — agents still sense, decide, and reason exactly as normal every tick, but the pheromone field never receives a deposit, so there is no stigmergic feedback channel at all. (Confirmed directly, not assumed: `ColonyTrialConfig.__post_init__` validates `decay`, `sensing_noise_std`, and `preference_variance`, but has no guard on `deposit_amount` — `0.0` is a legal, unvalidated configuration and constructs without error.)

## Experiment B — does the real mechanism beat a random-choice baseline? I

Hypothesis (Zero-deposit ablation closes the stigmergy-attribution confound)

*The real mechanism with pheromone deposit disabled (`deposit_amount=0.0` — the identical real `Agent/BeliefState/free-energy` decision loop as the full mechanism above, same seed sequence, same everything else) converges no better than the null model. Falsified by this zero-deposit condition's Wilson lower bound exceeding the null model's Wilson upper bound.*

**Real result** ( $N = 150$  trials, identical seed sequence `seed_base=0`, identical calibrated baseline configuration except `deposit_amount=0.0`):

## Experiment B — does the real mechanism beat a random-choice baseline? I

| harness                                     | successes/N | rate   | Wilson 95% CI    |
|---------------------------------------------|-------------|--------|------------------|
| full mechanism (deposit_amount=1.0)         | 140/150     | 0.9333 | (0.8816, 0.9634) |
| zero-deposit mechanism (deposit_amount=0.0) | 0/150       | 0.0000 | (0.0000, 0.0250) |
| null model (random choice)                  | 1/150       | 0.0067 | (0.0012, 0.0368) |

## Experiment B — does the real mechanism beat a random-choice baseline? I

$H_0$  **survives**: the zero-deposit condition's Wilson lower bound (0.0000) does not exceed the null model's Wilson upper bound (0.0368) — the two intervals overlap, and both are consistent with chance-level convergence. This is the confound-closing result, reported honestly regardless of which way it went (per this manuscript's own discipline): had the zero-deposit condition instead cleared the null model's upper bound, that would have been the surprising finding, requiring a different causal story (e.g. some residual bias in the free-energy decision rule unrelated to stigmergy) — it did not happen. Instead, severing only the pheromone deposit channel — while leaving every other part of the real decision loop untouched — collapses convergence all the way down to (numerically, even slightly below) the null model's own chance-level rate. A second, confirmatory comparison

## Experiment B — does the real mechanism beat a random-choice baseline? I

against the full mechanism makes the attribution explicit: the full mechanism's lower bound (0.8816) clears the zero-deposit condition's upper bound (0.0250) by a wide margin, and since `deposit_amount` is the *only* input changed between the two conditions (identical seed sequence, preferences, sensing-noise draws, decay, `num_agents`, `locations`, `num_ticks`), it is the single controlled variable responsible for the entire gap. Together, the two comparisons are the cleanest evidence in this manuscript that the real mechanism's advantage over chance is attributable specifically to the pheromone/ stigmergic feedback channel, not to some other property of the free-energy decision rule that would persist without it. **What this does and does not establish:** it establishes that removing the deposit channel alone is sufficient to erase the advantage in *this* configuration

## Experiment B — does the real mechanism beat a random-choice baseline? I

(num\_agents=8, two locations, num\_ticks=30, decay=0.46, sensing\_noise\_std=0.5); it does not establish that deposit is the *only* channel that could ever matter in some other configuration, nor does it characterize *how* the advantage degrades as deposit\_amount varies continuously between 0.0 and 1.0 — that dose-response question is untested here and is named as future work, not silently generalized past. Gated by test\_zero\_deposit\_real\_mechanism\_does\_not\_beat\_the\_null\_model\_closing\_the\_stigmergy\_confound and test\_zero\_deposit\_collapse\_relative\_to\_the\_full\_mechanism\_implicates\_the\_deposit\_channel\_specifically in test\_colony\_experiments\_extended.py.

Experiment C — does convergence rate decrease monotonically with heterogeneity magnitude? I

### Hypothesis (Monotonic heterogeneity sensitivity)

*Convergence rate is a monotonically non-increasing function of the width of preference\_mean\_range over four conditions — tight (9, 11), medium (8, 12) (the calibrated baseline), wide (5, 15), very-wide (2, 18) — holding every other parameter fixed. Falsified by any pair where a wider range scores a higher rate than a narrower one.*

Experiment C — does convergence rate decrease monotonically with heterogeneity magnitude? I

**Real result** (n=60 per condition, seed\_base=0):

Experiment C — does convergence rate decrease monotonically with heterogeneity magnitude? I

| condition | range   | width | successes/n | rate   | Wilson 95% CI    |
|-----------|---------|-------|-------------|--------|------------------|
| tight     | (9, 11) | 2.0   | 60/60       | 1.0000 | (0.9398, 1.0000) |
| medium    | (8, 12) | 4.0   | 56/60       | 0.9333 | (0.8407, 0.9738) |
| wide      | (5, 15) | 10.0  | 15/60       | 0.2500 | (0.1578, 0.3723) |
| very-wide | (2, 18) | 16.0  | 2/60        | 0.0333 | (0.0092, 0.1136) |

## Experiment C — does convergence rate decrease monotonically with heterogeneity magnitude? I

$H_0$  **survives** this experiment: the rate strictly decreases at every step of the sweep ( $1.0000 > 0.9333 > 0.2500 > 0.0333$ ), and the drop from medium to wide is the largest single step (nearly 70 percentage points) — heterogeneity does not degrade convergence gently across this range; it collapses sharply once agent preferences spread far enough that no single location can simultaneously satisfy enough of them. This is a genuine monotonic-decrease shape, not merely “the two extremes differ” — every intermediate step also holds the ordering. A single-statistic complement confirms the same shape as one number rather than three pairwise comparisons: the **Cochran–Armitage test for trend** (the same new colony/stats.py function used for the decay sweep) over the full ordered width set — successes  $\{60, 56, 15, 2\}$  out of 60 at widths  $\{2, 4, 10, 16\}$  — gives  $Z = -12.76$ ,

## Experiment C — does convergence rate decrease monotonically with heterogeneity magnitude? I

two-sided  $p < 10^{-10}$ : a strongly significant *decreasing* whole-sequence trend. Here — unlike the decay sweep, whose overall rise and local dip required reporting two tests — the trend statistic and the pairwise ordering agree in direction, so this is a confirmatory cross-check, not a second, distinct finding (gated by `test_cochran_armitage_finds_a_strong_negative_trend_a_cross_the_full_heterogeneity_sweep`).

**A previously uncomputed cross-check: does the mechanism's advantage over chance survive at the sweep's most extreme point?** Experiment C above establishes only a claim about *shape*: convergence decreases monotonically as `preference_mean_range` widens. It never asks the different question this cross-check asks: at the sweep's most extreme condition, very-wide (2, 18), is the real mechanism's convergence rate still statistically

## Experiment C — does convergence rate decrease monotonically with heterogeneity magnitude? I

distinguishable from Experiment B's random-choice null model, or has the mechanism's advantage over chance vanished by that point? Experiment B's null-model comparison was previously only ever run at the single calibrated medium (8, 12) baseline; it had never been crossed with any heterogeneity-sweep condition, very-wide included — grep-confirmed: no test or manuscript passage anywhere in this project instantiated `NullModelTrialConfig/run_null_model_trial` against any `preference_mean_range` other than that single baseline before this experiment. `NullModelTrialConfig` structurally has no `preference_mean_range` field at all (it has no preferences of any kind), so its convergence rate is entirely a function of `num_agents`, `locations`, and `num_ticks` — none of which the heterogeneity sweep varies — but it is still a

## Experiment C — does convergence rate decrease monotonically with heterogeneity magnitude? I

function of *seed*, so this experiment matches each real-mechanism seed base to a freshly-computed null-model run at the identical seed base, following the same “identical seed block” apples-to-apples principle Experiment B’s own zero-deposit ablation uses, rather than reusing a single null-model figure across different seed blocks.

### Hypothesis (Mechanism advantage over chance at the heterogeneity sweep’s extreme)

*The real mechanism’s convergence rate at the very-wide (2, 18) heterogeneity condition is no higher than the null-random-choice model’s rate at the identical `num_agents/locations/num_ticks` and the same seed block. Falsified (at a given seed base) by *very-wide*’s Wilson lower bound exceeding the null model’s Wilson upper bound AND a two-sided Fisher’s exact test on the two counts reaching  $p < 0.05$  at that seed base.*

Experiment C — does convergence rate decrease monotonically with heterogeneity magnitude? I

**Real result** ( $n=60$  for very-wide,  $N=150$  for the null model, each seed base's null-model run freshly computed at that seed base):

## Experiment C — does convergence rate decrease monotonically with heterogeneity magnitude? I

| seed base | very-wide<br>successes/n | rate   | Wilson 95% CI    | null successes/N | null Wilson 95%<br>CI | Fisher $p$ |
|-----------|--------------------------|--------|------------------|------------------|-----------------------|------------|
| 0         | 2/60                     | 0.0333 | (0.0092, 0.1136) | 1/150            | (0.0012, 0.0368)      | 0.1970     |
| 7000      | 5/60                     | 0.0833 | (0.0361, 0.1807) | 0/150            | (0.0000, 0.0250)      | 0.00168    |

## Experiment C — does convergence rate decrease monotonically with heterogeneity magnitude? I

$H_0$  gives a genuinely different answer at the two seed bases, and both are reported rather than one being selected as “the” result. At seed\_base=0,  $H_0$  **survives**: the Wilson intervals overlap substantially and the Fisher exact test does not clear  $p < 0.05$  ( $p = 0.1970$ ) — at this seed base, very-wide’s real-mechanism rate is **not** statistically distinguishable from random chance. At the disjoint seed\_base=7000,  $H_0$  **is rejected**: the exact test reaches  $p = 0.00168$  — at this seed base, the mechanism’s advantage over chance **does** survive at the same nominal condition. A scoping check on the next-most-heterogeneous condition, wide (5,15), confirms the ambiguity is specific to the sweep’s single most extreme point: wide clears its matching null-model baseline overwhelmingly at both seed bases (15/60 vs

Experiment C — does convergence rate decrease monotonically with heterogeneity magnitude? I

1/150,  $p = 2.14 \times 10^{-8}$  at seed\_base=0; 14/60 vs 0/150,  $p = 7.26 \times 10^{-9}$  at seed\_base=7000).

**What this does and does not establish.** It does not contradict Experiment C's monotonic-decrease claim about *shape* — that claim is about the ordering across widths within one seed base, and it replicates cleanly at both seed bases tested (Experiment C's own cross-seed replication, above). What this experiment adds is a different, previously unasked question — whether the mechanism's real-vs-chance advantage, established only at the calibrated baseline, survives at the sweep's most extreme point — and the honest answer, at only two seed bases tested, is “it depends which seed base you look at.” This is not resolved into a single verdict here: two seed bases are not enough to establish which result is “typical,” and no such

## Experiment C — does convergence rate decrease monotonically with heterogeneity magnitude? I

claim is made. It is reported as a genuine, previously-uncomputed, seed-base-dependent disagreement, directly relevant to this manuscript's central honesty discipline — a result that does not decompose cleanly is reported as such, not smoothed into whichever seed base tells a cleaner story. More seed bases, not a preference for one of the two already computed, is what would be needed to settle this further; that is named as future work, not silently resolved. Gated by `test_very_wide_at_seed0_does_not_clear_the_null_model_baseline`, `test_very_wide_at_seed7000_does_clear_the_null_model_baseline`, and `test_wide_condition_clears_the_null_model_baseline_at_both_seeds` in `test_colony_experiments_extended.py`.

## Honesty hedges common to all eleven analyses I

All eleven analyses hold `num_agents=8`, two locations, and `num_ticks=30` fixed — none of the results above is evidence about convergence at other colony sizes, other numbers of candidate locations, or other tick horizons, matching the scoping already stated for the `N=150` result. For the decay sweep, the cross-seed replication is likewise no longer merely a spot-check for one of its two previously-cited seed bases: an independent `seed_base=1000` run (`n=60` per value, the identical six decay values and calibrated baseline configuration as the `seed_base=0` sweep) is now a **gated** regression test (`test_decay_threshold_then_plateau_then_decline_shape_replicates_at_a_disjoint_seed_base`). At that disjoint block the successes are `0/60, 0/60, 58/60, 60/60, 60/60, 53/60` at decay `{0.10, 0.30, 0.46, 0.60, 0.80, 1.00}` —

## Honesty hedges common to all eleven analyses I

different exact counts from the `seed_base=0` run's 0/60, 0/60, 56/60, 60/60, 60/60, 56/60, as expected for a different seed block, but the identical qualitative shape: a near-zero floor at decay  $\leq 0.30$ , a 100% plateau at decay  $\in \{0.60, 0.80\}$ , and a measurable decline at decay = 1.00 relative to that plateau (53/60 here versus 60/60, compared to 56/60 versus 60/60 at `seed_base=0`). The exact successes counts and the exact magnitude of the top-end decline are not claimed to generalize beyond the pinned numbers at each seed base; it is the *shape* — threshold, then plateau, then decline — that replicates. The additional `seed_base=5000` spot-check remains an informal, ungated note (reproduced the same qualitative shape when checked by hand, but is not bound to a regression test). For the heterogeneity sweep, the cross-seed replication is likewise gated: an independent

## Honesty hedges common to all eleven analyses I

seed\_base=7000 run (a seed block sharing zero seeds with the seed\_base=0 sweep) is a **gated** regression test (test\_heterogeneity\_strict\_ordering\_replicates\_at\_a\_disjoint\_seed\_base). At that disjoint block the strict ordering reproduces with successes 60/54/14/5 (tight/medium/wide/very-wide) — different exact counts from the seed\_base=0 run's 60/56/15/2, as expected for a different seed block, but the identical monotonic-decrease shape. The exact successes counts at any given seed base are not claimed to generalize beyond the pinned numbers; it is the *ordering* that replicates. The mechanistic account offered for Experiment A was originally inference from trace inspection alone, explicitly flagged as distinct from the (better-supported) claim that decay has a real, non-monotonic causal effect on the rate within this controlled

## Honesty hedges common to all eleven analyses I

simulation; it is now additionally corroborated by a dedicated `sens ed_concentration_cap` ablation (above) showing that capping alone recovers convergence at  $\text{decay} \in \{0.10, 0.30\}$  — sufficiency evidence for the account, not a full mechanistic proof (see that section’s “What this does and does not establish”). Nothing here establishes these effects analytically (no closed-form model of convergence rate as a function of decay or heterogeneity width is derived or claimed) — every number in this section is an empirical measurement from a real, seeded simulation run, not a theoretical prediction.

**Why no multiple-comparisons correction (Bonferroni/Holm) is applied across these sweeps.** The decay sweep has six points and the heterogeneity sweep has four; eyeballing several points across a sweep for “the interesting difference” is exactly the kind of selective-inference exposure a multiple-comparisons correction exists to guard against, so it is worth stating plainly why none is applied here rather than leaving the omission implicit. First, each sweep’s *primary* whole-sequence claim rests on a single, pre-specified omnibus test — one Cochran–Armitage trend test per sweep (cochran\_armitage\_trend\_test,  $Z = +14.57$  for decay,  $Z = -12.76$  for heterogeneity, both  $p < 10^{-10}$ ) — not on scanning the sweep’s points pairwise for the smallest  $p$ -value. A single test per sweep has no multiple-comparisons problem to correct for. Second, the one genuinely exploratory *pairwise* comparison this manuscript

## Honesty hedges common to all eleven analyses I

does report and interpret — the decay  $\{0.60, 0.80\}$ -versus-1.00 dip, via `fisher_exact_test_two_sided` — is already disclosed above as **non-significant** ( $p = 0.1187$ , not clearing  $\alpha = 0.05$ ) and is explicitly not used to support a positive claim on its own: its evidentiary weight is stated to rest on the sweep's overall shape and its replication at disjoint seed bases, not on that one comparison clearing a significance threshold. A formal correction exists to prevent an uncorrected small  $p$ -value from a scan of many comparisons being over-interpreted as significant; that failure mode requires actually reporting an uncorrected pairwise  $p$ -value as if it, alone, established significance. This manuscript does not do that anywhere: the only pairwise test it reports is disclosed as failing to reach significance, and the only comparisons it treats as decisive are

## Honesty hedges common to all eleven analyses I

Wilson-interval-overlap checks between two pre-specified conditions (e.g. Experiment B's real-mechanism- versus-null and full-versus-zero-deposit comparisons), not a multiple-testing search over many candidate pairs for the one that clears a threshold. A correction is the right tool for a different failure mode than the one this manuscript's reporting practice actually exhibits.

The optional formal side-spec: shipped, not cut

## The optional formal side-spec: shipped, not cut I

Per ISC-35/36, this template ships both a Lean 4 model (`formal/lean/AntProtocol.lean`) and a TLA+ specification (`formal/tla/AntProtocol.tla + AntProtocol.cfg`) of the handshake protocol, each wired to exactly one real runnable check inside `scripts/check_formal_specs.sh` — `lake build` for the Lean model, and the TLC model checker (`tla2tools.jar`) for the TLA+ spec. Both checks are optional (they require `lake/elan` and a Java runtime respectively, and are not part of `--core-only` pipeline runs or the coverage gate on `src/`), but neither is decorative: the script fails non-zero if either check fails, and both have been executed at least once with captured output (recorded in `ISA.md`'s `## Decisions`). This is a deliberately different evidentiary register from the rest of this paper's claims — Lean/TLA+ model the

The optional formal side-spec: shipped, not cut I

*protocol design*, independent of whether this template's Python implementation actually conforms to that design; no claim in this paper conflates a passing Lean/TLA+ check with a proof about the Python code in `src/`.

Formal side-spec expansion: what grew, and what still needs wiring

Formal side-spec expansion: what grew, and what still needs wiring I

## Theorem (Mechanically verified protocol invariants)

*Three claims about the handshake protocol's design (not this template's Python implementation, per the scoping note above) are machine-checked, zero-sorry, zero-extra-axiom results, re-verified end to end by scripts/check\_formal\_specs.sh: (i) step\_to\_closed\_cases (Lean 4) — every transition into the closed phase originates from idle, established, or handshaking, never any other phase; (ii) cannot\_reuse\_consumed\_token (Lean 4) — a second, un-reassigned use call on a consumed session token has an unsatisfiable precondition; (iii) NoFalseEstablishment (TLA+, AntProtocolFaulty.tla) — a peer is never established without the other peer having genuinely sent the corresponding real message at some point (send-provenance), checked over 494 generated / 92 distinct states under arbitrary drop/corrupt/duplicate fault interleaving, and independently proven non-vacuous by a permanent negative control (below) rather than merely asserted.*

## Formal side-spec expansion: what grew, and what still needs wiring I

Both formal side-specs grew since the “shipped, not cut” decision above was first recorded, and both expansions were re-verified as real, runnable checks rather than taken on faith.

**A RedTeam pass found, and a mutation test confirmed, that clause (iii) above originally overclaimed what `NoFalseEstablishment` actually guards against.** The invariant is built on `iSentHello/rSentAck` — pure send-side history flags, set once by a genuine `SendHello`/accept-branch and never reset. That makes the invariant a check of **send-provenance** (“was a real message genuinely sent by the other peer, at some point”), not of **content-integrity resilience** (“was the specific delivered frame uncorrupted”). A direct mutation test made this concrete: widening `ReceiveHello`’s guard from `m.kind = "hello"` to also accept a “corrupt”-tagged frame enlarges the reachable

## Formal side-spec expansion: what grew, and what still needs wiring I

state space ( $92 \rightarrow 136$  distinct states) yet `NoFalseEstablishment` still reports zero violations — because `Corrupt` preserves a message's sender field and only ever acts on a frame a genuine `SendHello` already placed on the channel, so accepting a corrupted-but-genuinely-originated frame still satisfies send-provenance. This is exactly the same *near-vacuity* failure class already found and fixed twice on the Lean side (`no_direct_idle_to_established`, `closed_only_via_known_paths`) — an invariant that holds regardless of whether the mechanism it is supposed to gate is correctly implemented.

## Formal side-spec expansion: what grew, and what still needs wiring I

The fix mirrors this template's own tests/mypy\_fixtures/bad\_\*.py negative-control discipline rather than merely reworking the prose: formal/tla/AntProtocolFaultyNegControl.tla is a permanent, DELIBERATELY BROKEN sibling model — verbatim copy of AntProtocolFaulty.tla plus one added action, ForgeHello, which injects a [kind  $\rightarrow$  "hello", sender  $\rightarrow$  "I"] message onto the channel *without* setting iSentHello (an unauthenticated/spoofed hello the initiator never actually sent). scripts/check\_formal\_specs.sh runs TLC against it with **inverted** pass/fail logic — a reported violation is the expected, correct outcome; a clean “no error found” would mean the negative control itself is vacuous. Re-verified: TLC finds the violation in a 3-state trace (Init  $\rightarrow$  ForgeHello  $\rightarrow$

## Formal side-spec expansion: what grew, and what still needs wiring I

ReceiveHello establishes rPhase while iSentHello is still FALSE), proving NoFalseEstablishment genuinely depends on send-provenance holding, and clause (iii) above is now worded to the guarantee actually established rather than the broader one the original phrasing implied.

**Lean 4** (formal/lean/AntProtocol.lean) grew from 3 theorems to 7 audited property theorems (plus 3 reachability-witness theorems used as their preconditions, 10 declarations total), all zero-sorry, and — per the file's own #print axioms lines, re-confirmed by a clean lake build — depending on no axioms beyond Lean's own kernel: phase\_exhaustive (the four Phase constructors are the only inhabitants of the type — a background exhaustiveness fact the earlier revision left implicit), closed\_is\_terminal (no Step relation has

## Formal side-spec expansion: what grew, and what still needs wiring I

closed as its source — the direct Lean counterpart of “no outgoing edges from the ClosedSession phase” in `protocol/session.py`), a minimal affine `SessionToken/use` model plus `cannot_reuse_consumed_token` (the precondition a second, un-reassigned use call would need is unsatisfiable — the first Lean counterpart to the runtime `SessionConsumedError` discipline the type system cannot close, per sec. 3), and `no_direct_idle_to_established` — which *replaces* the prior `no_skip_to_established`, discarded because it reduced to a restatement of `established_requires_handshaking` via a vacuously-true `Reaches idle idle` conjunct rather than proving the one-step non-existence claim its name promised. Re-verified directly: `lake build` →

## Formal side-spec expansion: what grew, and what still needs wiring I

Build completed successfully (3 jobs), with all 7 #print axioms lines confirming zero extra axioms.

A second cross-vendor audit pass caught one more instance of the exact vacuity class `no_direct_idle_to_established` had already been introduced to fix: the theorem `closed_only_via_known_paths` concluded a three-way `Reaches idle established` `Reaches idle idle` `Reaches idle handshaking` disjunction, and `Reaches idle idle` is unconditionally provable via `Reaches.refl` regardless of which run was passed in — so the whole disjunction was satisfiable by a fixed proof term that never inspected its hypothesis. Rewording the conclusion to reference the run's own extracted predecessor phase did not escape the trap either, because in this small, 4-phase state machine every candidate predecessor is trivially reachable from `idle` by *some*

## Formal side-spec expansion: what grew, and what still needs wiring I

path independent of the specific run — the Reaches-wrapped framing is inherently satisfiable regardless of hypothesis in this model. The theorem was removed rather than patched a second time; the genuinely load-bearing, falsifiable content it was built on top of — `step_to_closed_cases`, a direct, universally-quantified claim about the `Step` relation itself (“every transition into `closed` comes from `idle`, `established`, or `handshaking` — never anything else”), false if a hypothetical sixth constructor targeted `closed` from elsewhere — was promoted to a named, `#print` axioms-checked theorem in its place. The audited theorem count stays at 7 (10 declarations including the 3 reachability witnesses); the set of what is proved changed to remove the vacuous member.

## Formal side-spec expansion: what grew, and what still needs wiring I

**TLA+** gained two things, not one. First, `AntProtocol.tla` added a liveness property, `HandshakeEventuallyResolves` ( $\Box(\text{handshaking} \Rightarrow \Diamond(\text{established} \vee \text{closed}))$ ), which only holds relative to the weak-fairness conjuncts (`WF_vars(Ack)`, `WF_vars(AbortHs)`) added to `Spec` — without them, TLC's exists-a-behavior semantics would permit a “coward” run that stalls in `handshaking` forever via stuttering, making the property vacuously unprovable. Re-verified: TLC's temporal-property check over the complete state space (5 distinct states) reports **“Model checking completed. No error has been found.”** Second, a genuinely new model shipped alongside the original: `AntProtocolFaulty.tla` (+ `.cfg`) is a two-peer (initiator/responder) model with an explicit message channel as

## Formal side-spec expansion: what grew, and what still needs wiring I

first-class state — not a copy of the single-peer `AntProtocol.tla`  
— with

`Send/ReceiveHello/ReceiveAck/Corrupt/DeliverCorrupt/Duplicate`  
actions mirroring `network/bus.py`'s real `drop/corrupt/duplicate`  
fault modes, checked against `NoFalseEstablishment` (a peer is  
never established without the other peer's real message having  
actually been sent, even under arbitrary fault interleaving — the  
direct TLA+ analogue of ISC-23/24). Re-verified: TLC reports  
**“Model checking completed. No error has been found”** over  
494 states generated, 92 distinct states.

Formal side-spec expansion: what grew, and what still needs wiring !

**Update — now wired.** `scripts/check_formal_specs.sh` was extended to run TLC a second time, against `AntProtocolFaulty.cfg/AntProtocolFaulty.tla`, immediately after the original `AntProtocol` check, and now a *third* time against the negative control `AntProtocolFaultyNegControl.cfg/.tla` described above (with inverted pass/fail logic) — all three TLC runs plus the Lean build must pass for the script to exit zero. Re-verified end to end: `lake build` (7 audited + 3 witness theorems, zero sorry) → PASS; TLC on `AntProtocol.tla` (liveness included) → PASS; TLC on `AntProtocolFaulty.tla` (494 states generated, 92 distinct) → PASS; TLC on `AntProtocolFaultyNegControl.tla` →

## Formal side-spec expansion: what grew, and what still needs wiring I

violation correctly detected in a 3-state trace → PASS; `check_for`  
`mal_specs.sh`: both formal side-specs passed, exit 0. All  
four checks now run under the same single automated command,  
closing the ISC-36 anti-criterion gap this section previously flagged  
as open — and, per the negative-control addition above, this  
section no longer merely asserts `NoFalseEstablishment` is  
non-vacuous, it demonstrates it the same way `tests/mypy_fixtu`  
`res/` demonstrates the mypy oracle is non-vacuous.

Threats to the claims, honestly stated

## Threats to the claims, honestly stated I

Nothing in this paper should be read as a claim that mypy --strict detects every illegal-state bug a determined adversary could construct in this codebase, nor that the fault-injection suite's four modes (drop/reorder/duplicate/corrupt) exhaust the space of network faults a production system would face — this template explicitly excludes real sockets, multi-process execution, and a message broker (Out of Scope). The colony integration test's stigmergic convergence result is a genuine, reproducible property of the specific three-agent, two-location, five-tick configuration tested; it is not evidence about convergence at colony scales, location counts, or tick horizons this template does not test. **The same caveat applies, separately, to the  $N = 150$  statistical claim above** (a RedTeam pass flagged that an earlier revision of this section scoped the caveat to the three-agent demo test only, leaving the headline

## Threats to the claims, honestly stated I

8-agent/2-location/30-tick statistical claim with no disclosed generalization limit of its own): “140/150 converge, Wilson lower bound 0.8816” is a property of *that* configuration, not a general law about stigmergic consensus at other agent counts, location counts, or tick horizons. Widening any of these claims would require the same discipline this paper asks of every claim already in it: a new negative-control test or fault-injected scenario before the wider claim is made, not merely a wider sentence in the manuscript. The pre-registered experiments above narrow this further, not widen it: they fix `num_agents=8`, two locations, and `num_ticks=30` throughout, and each experiment varies *exactly one* dimension (decay, mechanism-vs-random, heterogeneity width, or — for the zero-deposit ablation closing Experiment B’s stigmergy- attribution confound — `deposit_amount` alone)

## Threats to the claims, honestly stated I

while holding the others at the calibrated baseline. None of them licenses a claim about how decay, the mechanism's advantage over chance, heterogeneity, or the deposit-channel attribution would behave at a different colony size, location count, or tick horizon — only cross-seed robustness supports the narrower claim that the *qualitative shape* observed at `seed_base=0` is not an artifact of that one specific seed sequence, and that support is uneven by design: the heterogeneity ordering's replication at the disjoint `seed_base=7000` is now a gated regression test, whereas the decay threshold, the real-vs-null gap, and the zero-deposit ablation are each only established at `seed_base=0` (the decay threshold has additional informal spot-checks at other seed bases; the real-vs-null gap and the zero-deposit ablation do not, and neither is gated by a cross-seed-base regression test the way the heterogeneity ordering is).