Reproducibility: Seeded Regeneration and Artifact Trace

Reproducibility: Seeded Regeneration and Artifact Trace I

Re-running generation with seed 431 and the same lexicon produces the same token plan. The artifact set records `token_inventory.json`, `section_plan.json`, `injection_trace.json`, `manuscript_variables.json`, `figure_registry.json`, and the cover/results/configuration/evaluation figure set so the manuscript can be audited without reading the PDF.

The protocol emits MadlibConfig, review scenario, explicit/default path inventory, validated lexicon, digest input records, selection invariant set, TokenPlan, enabled section set, section variables, Markdown evidence tables, claim-aligned evidence surface, registered figure set, output/data, output/reports, and output/figures, output/manuscript, validated project output, review packet, copied publication-review bundle, fork migration notes. Project tests cover deterministic token choice, seed sensitivity, category-input sensitivity, malformed config

rejection, section disablement, artifact writing, and unresolved manuscript-token detection. The shared output validator then checks rendered PDFs, Markdown, figure registry, evidence registry, and design overlays.

The copied root output is therefore a consequence of local source and config. Generated files remain disposable; the durable contract is the ability to regenerate them from the tracked project tree and to observe the same validation gates passing.

Reproducibility: Seeded Regeneration and Artifact Trace I

- Config hash: `70a5c7317f5e44db`
- Generated: `2026-11-09T14:02:38Z` (derived from `SOURCE_DATE_EPOCH`; an explicit `not-recorded` marker is emitted when the reproducible timestamp is not supplied)
- Python: `3.14.10`
- Platform: `Linux x86_64`