

Reproducibility: Seeded Regeneration and Artifact Trace

## Reproducibility: Seeded Regeneration and Artifact Trace I

Re-running generation with seed 431 and the same lexicon produces the same token plan. The artifact set records `token_inventory.json`, `section_plan.json`, `injection_trace.json`, `manuscript_variables.json`, `figure_registry.json`, and the `cover/results/configuration/evaluation` figure set so the manuscript can be audited without reading the PDF.

The protocol emits `MadlibConfig`, review scenario, explicit/default path inventory, validated lexicon, digest input records, selection invariant set, `TokenPlan`, enabled section set, section variables, Markdown evidence tables, claim-aligned evidence surface, registered figure set, `output/data`, `output/reports`, and `output/figures`, `output/manuscript`, validated project output, review packet, copied publication-review bundle, fork migration notes. Project tests cover deterministic token choice, seed sensitivity, category-input sensitivity, malformed config

rejection, section disablement, artifact writing, and unresolved manuscript-token detection. The shared output validator then checks rendered PDFs, Markdown, figure registry, evidence registry, and design overlays.

The copied root output is therefore a consequence of local source and config. Generated files remain disposable; the durable contract is the ability to regenerate them from the tracked project tree and to observe the same validation gates passing.

## Reproducibility: Seeded Regeneration and Artifact Trace I

- ▶ Config hash: f64d6e672cc03a9f
- ▶ Generated: not-recorded (set SOURCE\_DATE\_EPOCH) (derived from SOURCE\_DATE\_EPOCH; an explicit not-recorded marker is emitted when the reproducible timestamp is not supplied)
- ▶ Python: 3.14.6
- ▶ Platform: Darwin arm64