

The AI Collaboration Model I

The Documentation Duality standard and three-tier skill architecture represent an empirical bet: that structured, machine-readable documentation measurably improves AI agent performance in research codebases. This section reports our key observations.

The documentation investment creates a positive feedback loop: as agents produce higher-quality outputs from structured context, developers maintain that documentation, which in turn improves future interactions (Lau and Guo 2025). We observed this concretely during `template/` development—each module's `SKILL.md` was refined through iterative AI-assisted generation, serving as both input prompt and output validator.

The AI Collaboration Model I

The SKILL.md layer, with its MCP-aligned YAML frontmatter (Anthropic 2024), provides a bridge to the agentic software paradigm. Lu et al.'s AI Scientist (Lu et al. 2024) demonstrates end-to-end autonomous research, while OpenHands is evaluated on SWE-Bench Verified (Wang et al. 2024; Jimenez et al. 2024). These systems motivate structured, protocol-aligned tool inventories for unfamiliar codebases. An agent navigating template/ reads CLAUDE.md for global constraints, scans AGENTS.md for local contracts, and can invoke skill-enabled capabilities through SKILL.md descriptors. The live module inventory verifies AGENTS.md and README.md coverage; SKILL.md remains a capability-specific layer rather than a universal claim about every infrastructure directory.

The AI Collaboration Model I

This three-tier model is, to our knowledge, novel in the research software engineering literature. The scale of the investment is substantial: 411 Markdown files under docs/ alone, plus an AGENTS.md/README.md pair in every directory and a SKILL.md descriptor on skill-enabled infrastructure modules. That count is itself injected from live introspection—the manuscript refuses to quote a documentation total it cannot recompute—and it represents a deliberate commitment to machine-readable context that shrinks the surface on which an agent can hallucinate.

The Learning Curve I

The Thin Orchestrator pattern imposes a cognitive overhead on researchers accustomed to writing monolithic scripts. The requirement to factor all logic into `src/` modules and use scripts only as stateless wiring introduces an additional layer of indirection. We mitigate this through:

The Learning Curve I

1. **Template exemplars:** `template_code_project` ships minimal optimization commentary; `template_prose_project` and `template_auresearch_project` broaden narrative + retrieval scaffolding.
2. **Documentation Duality:** Every directory has both `README.md` (for humans) and `AGENTS.md` (for AI collaborators), reducing the cost of navigation.
3. **Interactive orchestrator:** `run.sh` provides a TUI menu that abstracts pipeline complexity.
4. **Skill-level documentation:** The `docs/guides/` directory provides a progressive sequence of beginner, intermediate, advanced, and expert guides alongside a comprehensive new-project setup checklist.

Limitations I

Limitations I

- ▶ **LaTeX dependency:** The rendering pipeline requires a full TeX distribution (TeX Live or MiKTeX), which is a 4–6 GB install. This is the largest single dependency and is a barrier for researchers without system-level package management access.
- ▶ **Python-centric:** The infrastructure layer is Python-only. Projects in other languages can use the rendering and steganography stages but cannot leverage the `scientific` or `validation` modules.
- ▶ **Single-machine:** The pipeline runs locally. Distributed execution (e.g., across a compute cluster) is not natively supported, a gap where Snakemake, Nextflow, and CWL have clear superiority.
- ▶ **Steganographic robustness:** Alpha-channel overlays are stripped by aggressive PDF optimization tools (e.g., `qpdf --optimize`). QR codes are visible and removable. The current system provides *tamper detection* (via SHA-256 hashing) but not *non-repudiation* in the cryptographic sense—it lacks private-key digital signatures. An attacker with access to the source code could reproduce the watermark without having run the original pipeline.
- ▶ **Test duration:** The Zero-Mock policy increases test execution time from sub-second (mocked) to multi-minute (real) for the full infrastructure suite. This is acceptable for research workflows but may not suit continuous deployment scenarios.

Limitations II

- ▶ **AI-native writing tools:** `template/` does not include an integrated AI writing assistant comparable to Overleaf's Copilot features or OpenAI Prism's GPT-5.2 context-aware editing. The `infrastructure.llm` module provides LLM review as a pipeline stage but not as an interactive writing environment.