

Appendix: Documentation Inventory I

The repository maintains documentation at three levels:

Appendix: Documentation Inventory I

Table 1: Documentation inventory across the four-layer documentation architecture, from repository-wide system files to per-module skill descriptors.

Level	Files	Purpose
Repository root	AGENTS.md, CLAUDE.md, README.md, RUN_GUIDE.md	Global navigation and AI agent context
docs/ directory	411 files across 18 subdirectories	User guides, API reference, troubleshooting
Per-directory	AGENTS.md + README.md at every directory	Documentation Duality standard
Per-module (Tier 3)	SKILL.md at every infrastructure module	Machine-parseable MCP-aligned skill descriptor
Infrastructure-level (PAI)	PAI.md at infrastructure/ directory	Personal AI Infrastructure integration contract

Appendix: Documentation Inventory I

The docs/ subdirectories cover: core/ (essential docs), guides/ (progressive skill-level guides), architecture/ (system design), usage/ (content authoring), operational/ (build, config, logging, troubleshooting), reference/ (API, FAQ, glossary), modules/ (28 infrastructure modules), development/ (contributing, testing), maintenance/ (long-horizon maintenance: private projects, CI-local, regression, archival, bundle), best-practices/ (version control, migration), plans/ (strategic plans and decision records), rules/ (contributor norms), prompts/ (22 AI prompt templates), security/ (steganography, hashing), streams/ (livestream notes), and audit/ (review reports), plus the machine-generated _generated/ snippets and images/ assets.

Appendix: Documentation Inventory I

Every count in this appendix is injected from live repository introspection rather than hand-maintained: 411 counts every Markdown file beneath docs/ recursively, 18 counts its first-level subdirectories, and 22 counts the workflow subdirectories that each carry a SKILL.md descriptor. This is the same discipline the manuscript argues for throughout—a hand-typed documentation total silently rots as the tree grows, whereas a token re-resolves on every render. A reader onboarding to the repository should start at docs/core/, follow the graduated docs/guides/ skill ladder, and consult the per-directory AGENTS.md/README.md pair nearest to whatever code they are editing; AI agents additionally read each module's SKILL.md to locate capabilities without guessing API signatures.