# Appendix: Documentation Inventory I

The repository maintains documentation at three levels:

# Appendix: Documentation Inventory I

Table 1: Documentation inventory across the four-layer documentation architecture, from repository-wide system files to per-module skill descriptors.

| Level | Files | Purpose |
|---|---|---|
| Repository root | `AGENTS.md`, `.agents/`, `CLAUDE.md`, `README.md`, `docs/RUN_GUIDE.md` | Global navigation and AI agent context |
| `docs/` directory | 380 files across 20 subdirectories | User guides, API reference, troubleshooting |
| Per-directory | `AGENTS.md` + `README.md` in documented directories | Documentation Duality standard |
| Per-module (Tier 3) | `SKILL.md` on skill-enabled modules and public exemplars | Machine-parseable MCP-aligned skill descriptor |
| Infrastructure-level (PAI) | `PAI.md` at `infrastructure/` directory | Personal AI Infrastructure integration contract |

## Appendix: Documentation Inventory I

The `docs/` subdirectories cover: `core/` (essential docs), `guides/` (progressive skill-level guides), `architecture/` (system design), `usage/` (content authoring), `operational/` (build, config, logging, troubleshooting), `reference/` (API, FAQ, glossary), `modules/` (31 infrastructure modules), `development/` (contributing, testing), `maintenance/` (long-horizon maintenance: private projects, CI-local, regression, archival, bundle), `best-practices/` (version control, migration), `plans/` (strategic plans and decision records), `rules/` (contributor norms), `prompts/` (22 AI prompt templates), `security/` (steganography, hashing), `streams/` (livestream notes), and `audit/` (review reports), plus the machine-generated `_generated/` snippets and `images/` assets.

# Appendix: Documentation Inventory I

Every count in this appendix is injected from live repository introspection rather than hand-maintained: 380 counts every Markdown file beneath `docs/` recursively, 20 counts its first-level subdirectories, and 22 counts the workflow subdirectories that each carry a `SKILL.md` descriptor. This is the same discipline the manuscript argues for throughout—a hand-typed documentation total silently rots as the tree grows, whereas a token re-resolves on every render. A reader onboarding to the repository should start at `docs/core/`, follow the graduated `docs/guides/` skill ladder, and consult the per-directory `AGENTS.md`/`README.md` pair nearest to whatever code they are editing; AI agents begin with root `AGENTS.md` and `.agents/AGENTS.md`, then read local contracts and available `SKILL.md` descriptors. The shared skill manifest supplies

discovery; client-specific entrypoints reference the shared contracts.