Appendix: Documentation Inventory I

The repository maintains documentation at three levels:

Appendix: Documentation Inventory I

Table 1: Documentation inventory across the four-layer documentation architecture, from repository-wide system files to per-module skill descriptors.

Level	Files	Purpose
Repository root	`AGENTS.md`, `.agents/`, `CLAUDE.md`, `README.md`, `docs/RUN_GUIDE.md`	Global navigation and AI agent context
`docs/` directory	379 files across 20 subdirectories	User guides, API reference, troubleshooting
Per-directory	`AGENTS.md` + `README.md` in documented directories	Documentation Duality standard
Per-module (Tier 3)	`SKILL.md` on skill-enabled modules and public exemplars	Machine-parseable MCP-aligned skill descriptor
Infrastructure-level (PAI)	`PAI.md` at `infrastructure/` directory	Personal AI Infrastructure integration contract

Appendix: Documentation Inventory I

The `docs/` subdirectories cover: `core/` (essential docs), `guides/` (progressive skill-level guides), `architecture/` (system design), `usage/` (content authoring), `operational/` (build, config, logging, troubleshooting), `reference/` (API, FAQ, glossary), `modules/` (31 infrastructure modules), `development/` (contributing, testing), `maintenance/` (long-horizon maintenance: private projects, CI-local, regression, archival, bundle), `best-practices/` (version control, migration), `plans/` (strategic plans and decision records), `rules/` (contributor norms), `prompts/` (22 AI prompt templates), `security/` (steganography, hashing), `streams/` (livestream notes), and `audit/` (review reports), plus the machine-generated `_generated/` snippets and `images/` assets.

Appendix: Documentation Inventory I

Every count in this appendix is injected from live repository introspection rather than hand-maintained: 379 counts every Markdown file beneath `docs/` recursively, 20 counts its first-level subdirectories, and 22 counts the workflow subdirectories that each carry a `SKILL.md` descriptor. This is the same discipline the manuscript argues for throughout—a hand-typed documentation total silently rots as the tree grows, whereas a token re-resolves on every render. A reader onboarding to the repository should start at `docs/core/`, follow the graduated `docs/guides/` skill ladder, and consult the per-directory `AGENTS.md`/`README.md` pair nearest to whatever code they are editing; AI agents begin with root `AGENTS.md` and `.agents/AGENTS.md`, then read local contracts and available `SKILL.md` descriptors. The shared skill manifest supplies

discovery; client-specific entrypoints reference the shared contracts.